

# Video 19: pivoting tables with tidyr

Stats 102A

Miles Chen, PhD

# Pivoting Data

# Tidy Data

## The Philosophy of the Tidyverse

There are three rules which make a data set tidy:

- Every column is variable.
- Every row is an observation.
- Every cell is a single value.

# Tidy Data

storms

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ana	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arthur	45	1010	1996-06-21

# Tidy Data

storms

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ava	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arnur	40	1010	1996-06-21

- Storm name
- Wind Speed (mph)
- Air Pressure
- Date

We have one column for each variable

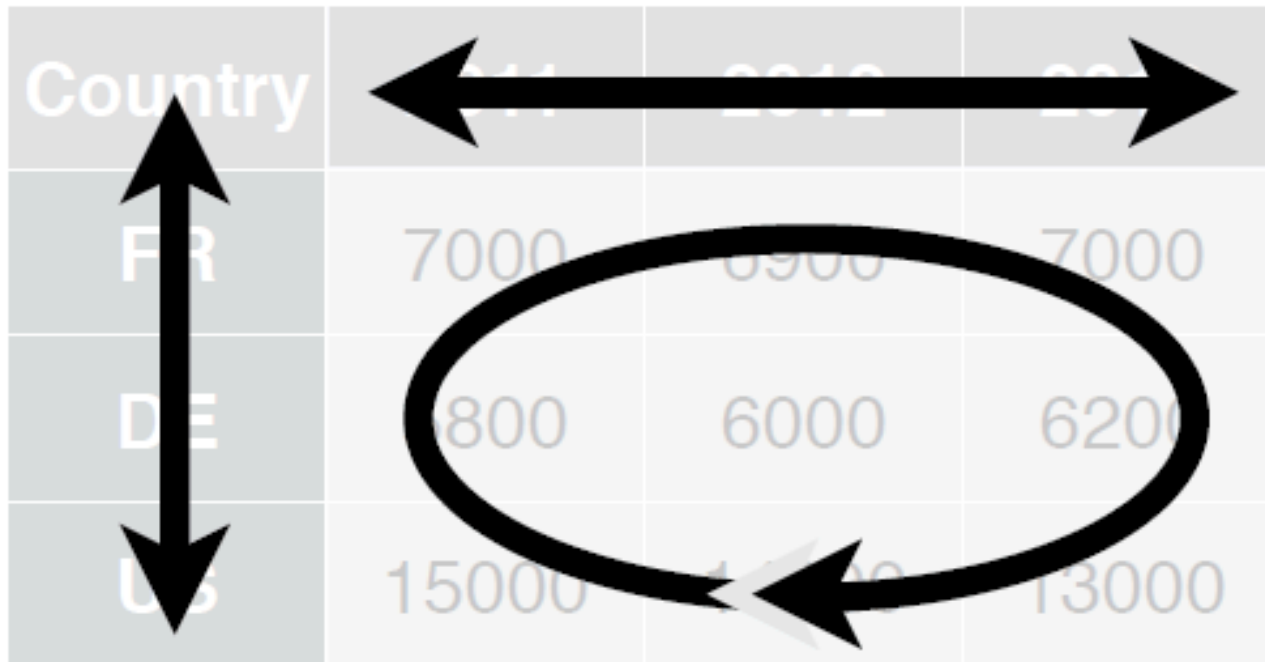
# Not Tidy Data

**cases**

Country	2011	2012	2013
FR	7000	6900	7000
DE	5800	6000	6200
US	15000	14000	13000

# Not Tidy Data

cases



The diagram illustrates a non-tidy data structure. It features a table with a 'Country' column and three columns representing years (2011, 2012, 2013). A vertical double-headed arrow on the left indicates that 'Country' is a single variable. A horizontal double-headed arrow at the top indicates that the years are spread across multiple columns. A large oval encircles the data values, indicating that the count of cases is spread across multiple columns.

Country	2011	2012	2013
FR	7000	6900	7000
DE	800	6000	6200
US	15000	14000	13000

- Country
- Year
- Count

One variable forms column headings, and the values are spread out across columns.

# Not Tidy Data

pollution

city	particle size	amount ( $\mu\text{g}/\text{m}^3$ )
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56



# Not Tidy Data

pollution

city	particle size	amount ( $\mu\text{g}/\text{m}^3$ )
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56

- City
- Amount of large particles
- Amount of small particles

The values of two different variables are stored in one column.

# Reading in the data

```
1 storms <- read_csv("https://raw.githubusercontent.com/rstudio/EDAWR/  
2 cases <- read_csv("https://raw.githubusercontent.com/rstudio/EDAWR/  
3 pollution <- read_csv("https://raw.githubusercontent.com/rstudio/ED
```

# Vectorized operations work for tidy data

```
1 storms$ratio <- storms$pressure / storms$wind
2 # 1007 / 110 = 9.15, 1009 / 45 = 22.4, etc.
3 storms
```

```
# A tibble: 6 × 5
```

	storm <chr>	wind <dbl>	pressure <dbl>	date <date>	ratio <dbl>
1	Alberto	110	1007	2000-08-03	9.15
2	Alex	45	1009	1998-07-27	22.4
3	Allison	65	1005	1995-06-03	15.5
4	Ana	40	1013	1997-06-30	25.3
5	Arlene	50	1010	1999-06-11	20.2
6	Arthur	45	1010	1996-06-17	22.4

# The Cases table

1 cases

```
# A tibble: 3 × 4
```

	country	`2011`	`2012`	`2013`
	<chr>	<dbl>	<dbl>	<dbl>
1	FR	7000	6900	7000
2	DE	5800	6000	6200
3	US	15000	14000	13000

# Getting things tidy

The variables are: country, year, count


If we want to make this tidy, what will be the dimensions of the resulting data?

**cases**

Country	2011	2012	2013
FR	7000	6900	7000
DE	5800	6000	6200
US	15000	14000	13000

# Result: a 9 x 3 tibble

Country	2011	2012	2013
FR	7000	6900	7000
DE	5800	6000	6200
US	15000	14000	13000



Country	Year	n
FR	2011	7000
DE	2011	5800
US	2011	15000
FR	2012	6900
DE	2012	6000
US	2012	14000
FR	2013	7000
DE	2013	6200
US	2013	13000

# pivot\_longer()

To achieve the desired result, we use the function `pivot_longer()` because we want the resulting data set to be longer than the original data. (older versions called this `gather()`)

```
1 pivot_longer(cases,  
2             cols = "2011":"2013",  
3             names_to = "year",  
4             values_to = "cases")
```

```
# A tibble: 9 × 3  
  country year  cases  
  <chr>   <chr> <dbl>  
1 FR     2011    7000  
2 FR     2012    6900  
3 FR     2013    7000  
4 DE     2011    5800  
5 DE     2012    6000  
6 DE     2013    6200  
7 US     2011   15000  
8 US     2012   14000  
9 US     2013   13000
```

# The `pivot_longer()` function

```
1 pivot_longer(data = cases,  
2             cols = "2011":"2013",  
3             names_to = "year",  
4             values_to = "cases")
```

The `pivot_longer()` function takes in a few arguments:

- `data` is the name of the data.frame or tibble that we will pivot
- `cols` are the names of the columns that will be pivoted. In this case, we want the columns named “2011” through “2013”. With `tidyr`, you can specify a range of column names with the `:` operator. Otherwise, you can provide a vector of column names
- `names_to` is a character string with what you want to call the resulting column of names. The former column names will be put into this column.
- `values_to` is a character string with what you want to call the resulting column of values. The former cell values will be put into this column.



# The names are arbitrary

```
1 pivot_longer(cases,  
2               cols = "2011":"2013",  
3               names_to = "when it happened",  
4               values_to = "how many")
```

```
# A tibble: 9 × 3
```

	country	`when it happened`	`how many`
	<chr>	<chr>	<dbl>
1	FR	2011	7000
2	FR	2012	6900
3	FR	2013	7000
4	DE	2011	5800
5	DE	2012	6000
6	DE	2013	6200
7	US	2011	15000
8	US	2012	14000
9	US	2013	13000

# What happens if?

```
# A tibble: 3 × 4
  country `2011` `2012` `2013`
  <chr>    <dbl>  <dbl>  <dbl>
1 FR      7000    6900   7000
2 DE      5800    6000   6200
3 US     15000   14000  13000
```

What happens if the columns I pivot are only “2012” and “2013”?

```
1 pivot_longer(cases,
2               cols = "2012":"2013",
3               names_to = "year",
4               values_to = "cases")
```

# Answer

What happens if the columns I pivot are only “2012” and “2013”?

The columns that are not pivoted are duplicated for the new rows created.

```
1 pivot_longer(cases,  
2               cols = "2012":"2013",  
3               names_to = "year",  
4               values_to = "cases")
```

```
# A tibble: 6 × 4
```

	country	`2011`	year	cases
	<chr>	<dbl>	<chr>	<dbl>
1	FR	7000	2012	6900
2	FR	7000	2013	7000
3	DE	5800	2012	6000
4	DE	5800	2013	6200
5	US	15000	2012	14000
6	US	15000	2013	13000

# What happens if?

```
# A tibble: 3 × 4
  country `2011` `2012` `2013`
  <chr>    <dbl>  <dbl>  <dbl>
1 FR      7000    6900   7000
2 DE      5800    6000   6200
3 US     15000   14000  13000
```

What happens if I include “country” in the columns I pivot?

```
1 pivot_longer(cases,
2               cols = "country":"2013",
3               names_to = "name",
4               values_to = "value")
```

# Answer

What happens if I include “country” in the columns I pivot?

```
1 pivot_longer(cases,  
2             cols = "country":"2013",  
3             names_to = "name",  
4             values_to = "value")
```

```
Error in `pivot_longer()`:  
! Can't combine `country` <character> and `2011` <double>.
```

# Answer

What happens if I include “country” in the columns I pivot?  
(I’ve converted everything to character.)

```
1 cases2 <- data.frame(lapply(cases[,1:4], as.character))
2 names(cases2) <- c("country", "2011", "2012", "2013")
3 pivot_longer(cases2,
4               cols = "country":"2013",
5               names_to = "name",
6               values_to = "value")
```

# A tibble: 12 × 2

	name	value
	<chr>	<chr>
1	country	FR
2	2011	7000
3	2012	6900
4	2013	7000
5	country	DE
6	2011	5800
7	2012	6000
8	2013	6200
9	country	US
10	2011	15000

11	2012	14000
12	2013	13000

# The pollution table

```
1 pollution
```

```
# A tibble: 6 × 3
```

	city	size	amount
	<chr>	<chr>	<dbl>
1	New York	large	23
2	New York	small	14
3	London	large	22
4	London	small	16
5	Beijing	large	121
6	Beijing	small	56



# Getting things tidy

The variables are: city, large particle amount, small particle amount

If we want to make this tidy, what will be the dimensions of the resulting data?

# Result: a 3 x 3 tibble

city	size	amount
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56



city	large	small
New York	23	14
London	22	16
Beijing	121	56

# `pivot_wider()`

To achieve the desired result, we use the function `pivot_wider()` because we want the resulting dataset to be wider than the original data. (older versions called this `spread()`)

```
1 pivot_wider(pollution,  
2             names_from = "size",  
3             values_from = "amount")
```

```
# A tibble: 3 × 3
```

	city	large	small
	<chr>	<dbl>	<dbl>
1	New York	23	14
2	London	22	16
3	Beijing	121	56

# `pivot_wider()`

```
1 pivot_wider(pollution,  
2             names_from = "size",  
3             values_from = "amount")
```

The `pivot_wider()` function takes in a few arguments:

- `data` is the name of the data.frame or tibble that we will pivot
- `names_from` is the name of the column that has the names that will become column headers
- `values_from` is the name of the column that has the values

# `pivot_wider()` is sensitive to spelling differences

What will happen?

```
1 pollution2 <- pollution
2 pollution2[1,1] <- "NYC"
3 pollution2
```

```
# A tibble: 6 × 3
  city      size amount
<chr>    <chr>   <dbl>
1 NYC      large     23
2 New York small     14
3 London   large     22
4 London   small     16
5 Beijing  large    121
6 Beijing  small     56
```

```
1 pivot_wider(pollution2,
2             names_from = "size",
3             values_from = "amount")
```

# Result

```
1 pivot_wider(pollution2,  
2             names_from = "size",  
3             values_from = "amount")
```

# A tibble: 4 × 3

	city <chr>	large <dbl>	small <dbl>
1	NYC	23	NA
2	New York	NA	14
3	London	22	16
4	Beijing	121	56

# Result

Sometimes you truly do have a scenario where you want to pivot wider and some entries do not exist. If you don't want **NAs** to show, you can specify a fill value.

```
1 pivot_wider(pollution2,  
2             names_from = "size",  
3             values_from = "amount",  
4             values_fill = 0)
```

```
# A tibble: 4 × 3  
  city      large small  
  <chr>    <dbl> <dbl>  
1 NYC          23     0  
2 New York     0     14  
3 London       22     16  
4 Beijing     121     56
```

# Another example

What will happen?

```
1 pollution2 <- pollution
2 pollution2[1,2] <- "LARGE"
3 pollution2
```

# A tibble: 6 × 3

	city	size	amount
	<chr>	<chr>	<dbl>
1	New York	LARGE	23
2	New York	small	14
3	London	large	22
4	London	small	16
5	Beijing	large	121
6	Beijing	small	56

```
1 pivot_wider(pollution2,
2             names_from = "size",
3             values_from = "amount")
```



# Result

```
1 pivot_wider(pollution2,  
2             names_from = "size",  
3             values_from = "amount")
```

# A tibble: 3 × 4

	city	LARGE	small	large
	<chr>	<dbl>	<dbl>	<dbl>
1	New York	23	14	NA
2	London	NA	16	22
3	Beijing	NA	56	121

# `pivot_longer()` and `pivot_wider()` are inverse operations

```
1 pollution
```

```
# A tibble: 6 × 3  
  city      size amount  
  <chr>    <chr> <dbl>  
1 New York large    23  
2 New York small    14  
3 London   large    22  
4 London   small    16  
5 Beijing  large   121  
6 Beijing  small    56
```

```
1 w <- pivot_wider(pollution, names_from = "size", values_from = "amount")  
2 w
```

```
# A tibble: 3 × 3  
  city      large small  
  <chr>    <dbl> <dbl>  
1 New York    23    14  
2 London     22    16  
3 Beijing   121    56
```

# `pivot_longer()` and `pivot_wider()` are inverse operations

```
1 w
```

```
# A tibble: 3 × 3  
  city      large small  
  <chr>    <dbl> <dbl>  
1 New York      23     14  
2 London        22     16  
3 Beijing     121     56
```

```
1 pivot_longer(w, cols = "large":"small", names_to = "size", values_to = "amount")
```

```
# A tibble: 6 × 3  
  city      size amount  
  <chr>    <chr>  <dbl>  
1 New York large      23  
2 New York small      14  
3 London  large      22  
4 London  small      16  
5 Beijing large     121  
6 Beijing small      56
```

# `pivot_longer()` and `pivot_wider()` are inverse operations

```
1 cases
```

```
# A tibble: 3 × 4
  country `2011` `2012` `2013`
  <chr>    <dbl>  <dbl>  <dbl>
1 FR      7000    6900    7000
2 DE      5800    6000    6200
3 US     15000   14000   13000
```

```
1 1 <- pivot_longer(cases, cols = "2011":"2013", names_to = "year", values_to = "count")
2 1
```

```
# A tibble: 9 × 3
  country year  count
  <chr>   <chr> <dbl>
1 FR    2011    7000
2 FR    2012    6900
3 FR    2013    7000
4 DE    2011    5800
5 DE    2012    6000
6 DE    2013    6200
7 US    2011   15000
8 US    2012   14000
9 US    2013   13000
```

# `pivot_longer()` and `pivot_wider()` are inverse operations

```
1 1
```

```
# A tibble: 9 × 3
  country year  count
  <chr>   <chr> <dbl>
1 FR     2011    7000
2 FR     2012    6900
3 FR     2013    7000
4 DE     2011    5800
5 DE     2012    6000
6 DE     2013    6200
7 US     2011   15000
8 US     2012   14000
9 US     2013   13000
```

```
1 pivot_wider(1, names_from = "year", values_from = "count")
```

```
# A tibble: 3 × 4
  country `2011` `2012` `2013`
  <chr>   <dbl> <dbl> <dbl>
1 FR      7000   6900   7000
2 DE      5800   6000   6200
3 US     15000  14000  13000
```