

Video 20: dplyr - select, filter, mutate

Stats 102A

Miles Chen, PhD

dplyr

dplyr

`dplyr` is a core part of the tidyverse.

You can load the library with `library(dplyr)` or by loading all of the tidyverse with `library(tidyverse)`

dplyr vignette

When working with data you must:

- Figure out what you want to do.
- Describe those tasks in the form of a computer program.
- Execute the program.

The dplyr package makes these steps fast and easy:

- By constraining your options, it helps you think about your data manipulation challenges.
- It provides simple “verbs”, functions that correspond to the most common data manipulation tasks, to help you translate your thoughts into code.

dplyr vignette

dplyr is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges:

- `select()` picks variables based on their names.
- `filter()` picks cases based on their values.
- `mutate()` adds new variables that are functions of existing variables.
- `arrange()` changes the ordering of the rows.
- `summarise()` reduces multiple values down to a single summary.

These all combine naturally with `group_by()` which allows you to perform any operation “by group.”

the dplyr cheat sheet

<https://github.com/rstudio/cheatsheets/blob/master/data-transformation.pdf>

Data Transformation with dplyr : CHEAT SHEET

dplyr

dplyr functions work with pipes and expect tidy data. In tidy data:

Each variable is in its own column. Each observation, or case, is in its own row. $x \%>\% f(y)$ becomes $f(x, y)$

Summarise Cases

These apply **summary functions** to columns to create a new table of summary statistics. Summary functions take vectors as input and return one value (see back).

summary function

```
summarise(data, ...)
```

Compute table of summaries.
`summarise(mtcars, avg = mean(mpg))`

```
count(x, ...)
```

Count number of rows in each group defined by the variables in ... Also `tally()`.
`count(mtcars, Species)`

VARIATIONS

summarise_all() - Apply funs to every column.
summarise_at() - Apply funs to specific columns.
summarise_if() - Apply funs to all cols of one type.

Group Cases

Use **group_by()** to create a "grouped" copy of a table. dplyr functions will manipulate each "group" separately and then combine the results.

group_by(data, ...) add = FALSE
Returns copy of table grouped by ...
`g_mtcars <- group_by(mtcars, Species)`

ungroup(x, ...)
Returns ungrouped copy of table.
`ungroup(g_mtcars)`

Manipulate Cases

EXTRACT CASES

Row functions return a subset of rows as a new table.

filter(data, ...) Extract rows that meet logical criteria. `filter(mtcars, Sepal.Length > 7)`

distinct(data, ...) Remove rows with duplicate values. `distinct(mtcars, Species)`

sample_frac(tbl, size = 1, replace = FALSE, weight = NULL, env = parent.frame()) Randomly select fraction of rows.
`sample_frac(mtcars, 0.5, replace = TRUE)`

sample_n(tbl, size, replace = FALSE, weight = NULL, env = parent.frame()) Randomly select size rows. `sample_n(mtcars, 10, replace = TRUE)`

slice(data, ...) Select rows by position.
`slice(mtcars, 10:15)`

top_n(x, n, wt) Select and order top n entries (by group if grouped data). `top_n(mtcars, 5, Sepal.Width)`

Logical and boolean operators to use with filter()

<	<=	is.na()	%in%	is.na()	!	&	xor()
>	>=	is.na()	!	&	xor()		

See **Base::logic** and **TComparison** for help.

ARRANGE CASES

arrange(data, ...) Order rows by values of a column or columns (low to high), use with **desc()** to order from high to low.
`arrange(mtcars, desc(mpg))`

ADD CASES

add_row(data, ...) before = NULL, after = NULL
Add one or more rows to a table.
`add_row(mtcars, new = 1:32)`

Manipulate Variables

EXTRACT VARIABLES

Column functions return a set of columns as a new vector or table.

pull(data, var = 1) Extract column values as a vector. Choose by name or index.
`pull(mtcars, Sepal.Length)`

select(data, ...) Extract columns as a table. Also **select_if()**.
`select(mtcars, Sepal.Length, Species)`

Use these helpers with **select()**, e.g. `select(mtcars, starts_with("Sepal"))`

contains(match) - one of ...
ends_with(match) - one of ...
matches(match) - starts with (match)

MAKE NEW VARIABLES

These apply **vectorized functions** to columns. Vectorized funs take vectors as input and return vectors of the same length as output (see back).

mutate(data, ...) Compute new column(s).
`mutate(mtcars, gpm = 1/mpg)`

transmute(data, ...) Compute new column(s), drop others.
`transmute(mtcars, gpm = 1/mpg)`

mutate_all(tbl, funs, ...) Apply funs to every column. Use with **funs()**. Also **mutate_if()**.
`mutate_all(mtcars, funs(log))`

mutate_at(tbl, cols, funs, ...) Apply funs to specific columns. Use with **funs()**, **vars()** and the helper functions for **select()**.
`mutate_at(mtcars, vars(Sepal.Length), funs(log))`

add_column(data, ...) before = NULL, after = NULL
Add new column(s). Also **add_count()**.
`add_tally(mtcars, new = 1:32)`

rename(mtcars, Length = Sepal.Length)

Vector Functions

TO USE WITH MUTATE()

mutate() and **transmute()** apply vectorized functions to columns to create new columns. Vectorized functions take vectors as input and return vectors of the same length as output.

vectorized function

OFFSETS
`dplyr::lag()` - Offset elements by 1
`dplyr::lead()` - Offset elements by -1

CUMULATIVE AGGREGATES

`dplyr::cumall()` - Cumulative all()
`dplyr::cumany()` - Cumulative any()
`dplyr::cummax()` - Cumulative max()
`dplyr::cummean()` - Cumulative mean()
`dplyr::cummin()` - Cumulative min()
`dplyr::cumprod()` - Cumulative prod()
`dplyr::cumsum()` - Cumulative sum()

RANKINGS

`dplyr::cume_dist()` - Proportion of all values <= value
`dplyr::dense_rank()` - rank with ties = min, no gaps
`dplyr::min_rank()` - rank with ties = min
`dplyr::rank()` - rank scaled to [0,1]
`dplyr::row_number()` - rank with ties = "first"

MATH

`+`, `-`, `*`, `/`, `%`, `%%`, `%/%` - arithmetic ops
`log()`, `log2()`, `log10()` - logs
`<`, `<=`, `>`, `>=`, `==`, `!=`, `is.na()` - logical comparisons
`dplyr::between()` - `x >= left & x <= right`
`dplyr::near()` - safe == for floating point numbers

MISC

`dplyr::case_when()` - multi-case if, else
`dplyr::coalesce()` - first non-NA values by element across a set of vectors
`dplyr::if_else()` - element-wise if, else
`dplyr::na_if()` - replace specific values with NA
`dplyr::pmax()` - element-wise max()
`dplyr::pmin()` - element-wise min()
`dplyr::recode()` - Vectorized switch()
`dplyr::recode_factor()` - Vectorized switch() for factors

Summary Functions

TO USE WITH SUMMARISE()

Summary functions apply summary functions to columns to create a new table. Summary functions take vectors as input and return single values as output.

summary function

COUNTS
`dplyr::n()` - number of values/rows
`dplyr::n_distinct()` - # of unique values
`sum(is.na())` - # of non-NA's

LOCATION

`mean()` - mean, also `mean(is.na())`
`median()` - median

LOGICALS

`mean()` - Proportion of TRUE's
`sum()` - # of TRUE's

POSITION/ORDER

`dplyr::first()` - first value
`dplyr::last()` - last value
`dplyr::nth()` - value in nth location of vector

RANK

`quantile()` - nth quantile
`min()` - minimum value
`max()` - maximum value

SPREAD

`IQR()` - Inter-Quartile Range
`mad()` - median absolute deviation
`sd()` - standard deviation
`var()` - variance

Row Names

Tidy data does not use rownames, which store a variable outside of the columns. To work with the rownames, first move them into a column.

rownames_to_column()
Move rownames to column.
`rownames_to_column(mtcars, var = "C")`

column_to_rownames()
Move col in row names.
`column_to_rownames(mtcars, col = "C")`

Also has **rownames()**, **remove_rownames()**

Combine Tables

COMBINE VARIABLES

Use **bind_cols()** to paste tables beside each other as they are.

bind_cols(...) Returns tables placed side by side as a single table. BE SURE THAT ROWS ALIGN.

Use a "Mutating Join" to join one table to columns from another, matching values with the rows that they correspond to. Each join retains a different combination of values from the tables.

left_join(x, y, by = NULL, copy = FALSE, suffixes = c("x", "y"), ...)
Join matching values from y to x.

right_join(x, y, by = NULL, copy = FALSE, suffixes = c("x", "y"), ...)
Join matching values from x to y.

inner_join(x, y, by = NULL, copy = FALSE, suffixes = c("x", "y"), ...)
Join data. Retain only rows with matches.

full_join(x, y, by = NULL, copy = FALSE, suffixes = c("x", "y"), ...)
Join data. Retain all values, all rows.

Use **by = c("col1", "col2", ...)** to specify one or more common columns to match on.
`left_join(x, y, by = "X")`

Use a named vector, **by = c("col1" = "col2", ...)**, to match on columns that have different names in each table.
`left_join(x, y, by = c("C" = "D"))`

Use **suffix** to specify the suffix to give to unmatched columns that have the same name in both tables.
`left_join(x, y, by = c("C" = "D"), suffix = c("1", "2"))`

Use a "Filtering Join" to filter one table against the rows of another.

semi_join(x, y, by = NULL, ...)
Return rows of x that have a match in y. USEFUL TO SEE WHAT WILL BE JOINED.

anti_join(x, y, by = NULL, ...)
Return rows of x that do not have a match in y. USEFUL TO SEE WHAT WILL NOT BE JOINED.

COMBINE CASES

Use **bind_rows()** to paste tables below each other as they are.

bind_rows(..., id = NULL)
Returns tables one on top of the other as a single table. Set `id` to a column name to add a column of the original table names (as pictured).

intersect(x, y, ...)
Rows that appear in both x and y.

setdiff(x, y, ...)
Rows that appear in x but not y.

union(x, y, ...)
Rows that appear in x or y. (Duplicates removed, union_all) retains duplicates.

Use **setequal()** to test whether two data sets contain the exact same rows (in any order).

EXTRACT ROWS

slice()

R Studio

RStudio is a trademark of RStudio, Inc. • CC BY SA RStudio • info@rstudio.com • 844-448-1212 • rstudio.com • Learn more with `browser(gettext(package = "dplyr", "RStudio"))` • dplyr 0.7.0 • 10066 • 2.2.0 • Updated 2017-03

R Studio

RStudio is a trademark of RStudio, Inc. • CC BY SA RStudio • info@rstudio.com • 844-448-1212 • rstudio.com • Learn more with `browser(gettext(package = "dplyr", "RStudio"))` • dplyr 0.7.0 • 10066 • 2.2.0 • Updated 2017-03

dplyr examples

The `starwars` data set

The Star Wars data set is included with `dp1yr`. It contains information about various Star Wars characters from the first 7 Star Wars movies.

```
1 starwars
```

```
# A tibble: 87 × 14
```

	name	height	mass	hair_color	skin_color	eye_color	birth_year	sex
gender								
	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>
<chr>								
1	Luke Sk...	172	77	blond	fair	blue	19	male
mascu...								
2	C-3PO	167	75	<NA>	gold	yellow	112	none
mascu...								
3	R2-D2	96	32	<NA>	white, bl...	red	33	none
mascu...								
4	Darth V...	202	136	none	white	yellow	41.9	male
mascu...								
5	Leia Or...	150	49	brown	light	brown	19	fema...
femin...								
6	Obi-wan	178	136	brown	light	blue	50	male

Select columns with `select()`

When using `select()`, you do not need to put quotes around the column names if there are no spaces in the names.

```
1 select(starwars, name, homeworld, species, films)
```

```
# A tibble: 87 × 4
```

	name	homeworld	species	films
	<chr>	<chr>	<chr>	<list>
1	Luke Skywalker	Tatooine	Human	<chr [5]>
2	C-3PO	Tatooine	Droid	<chr [6]>
3	R2-D2	Naboo	Droid	<chr [7]>
4	Darth Vader	Tatooine	Human	<chr [4]>
5	Leia Organa	Alderaan	Human	<chr [5]>
6	Owen Lars	Tatooine	Human	<chr [3]>
7	Beru Whitesun Lars	Tatooine	Human	<chr [3]>
8	R5-D4	Tatooine	Droid	<chr [1]>
9	Biggs Darklighter	Tatooine	Human	<chr [1]>
10	Obi-Wan Kenobi	Stewjon	Human	<chr [6]>

```
# i 77 more rows
```

Using the pipe

The pipe `%>%` takes the result of what is in front of the pipe and inserts it as the first argument in the function that comes after the pipe. `x %>% f(y)` turns into `f(x, y)` so the result from one step is then “piped” into the next step.

```
1 # select(starwars, name, homeworld, species, films) is exactly equivalent to
2 starwars %>%
3   select(name, homeworld, species, films)
```

```
# A tibble: 87 × 4
```

	name <chr>	homeworld <chr>	species <chr>	films <list>
1	Luke Skywalker	Tatooine	Human	<chr [5]>
2	C-3PO	Tatooine	Droid	<chr [6]>
3	R2-D2	Naboo	Droid	<chr [7]>
4	Darth Vader	Tatooine	Human	<chr [4]>
5	Leia Organa	Alderaan	Human	<chr [5]>
6	Owen Lars	Tatooine	Human	<chr [3]>
7	Beru Whitesun Lars	Tatooine	Human	<chr [3]>
8	R5-D4	Tatooine	Droid	<chr [1]>
9	Biggs Darklighter	Tatooine	Human	<chr [1]>
10	Obi-Wan Kenobi	Stewjon	Human	<chr [6]>

```
# i 77 more rows
```

Native Pipe | >

R versions 4.1 and later have a native pipe. | >

The functionality is almost identical to the pipe that is part of the tidyverse %>%

More information: <https://www.tidyverse.org/blog/2023/04/base-vs-magrittr-pipe/>

```
1 # select(starwars, name, homeworld, species, films) is exactly equivalent to
2 starwars |>
3   select(name, homeworld, species, films)
```

```
# A tibble: 87 × 4
  name                homeworld species films
  <chr>              <chr>    <chr>  <list>
1 Luke Skywalker    Tatooine Human  <chr [5]>
2 C-3PO             Tatooine Droid  <chr [6]>
3 R2-D2             Naboo    Droid  <chr [7]>
4 Darth Vader       Tatooine Human  <chr [4]>
5 Leia Organa       Alderaan Human  <chr [5]>
6 Owen Lars         Tatooine Human  <chr [3]>
7 Beru Whitesun Lars Tatooine Human  <chr [3]>
8 R5-D4             Tatooine Droid  <chr [1]>
9 Biggs Darklighter Tatooine Human  <chr [1]>
10 Obi-Wan Kenobi    Stewjon  Human  <chr [6]>
# i 77 more rows
```

Shortcut to insert the pipe

Shortcut to insert the pipe:

CTRL(CMD) + SHIFT + M

Select columns with `select()`

- Use a negative sign to deselect columns

```
1 starwars %>%  
2   select( -name, -eye_color, -birth_year) %>%  
3   head(3)
```

A tibble: 3 × 11

	height	mass	hair_color	skin_color	sex	gender	homeworld	species	films
	<int>	<dbl>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<list>
1	172	77	blond	fair	male	masculine	Tatooine	Human	<chr>
2	167	75	<NA>	gold	none	masculine	Tatooine	Droid	<chr>
3	96	32	<NA>	white, blue	none	masculine	Naboo	Droid	<chr>

i 2 more variables: vehicles <list>, starships <list>

select() example

- Use colon notation to select a range of columns

```
1 starwars %>%  
2   select(name:eye_color) %>%  
3   head(3)
```

A tibble: 3 × 6

	name	height	mass	hair_color	skin_color	eye_color
	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>
1	Luke Skywalker	172	77	blond	fair	blue
2	C-3PO	167	75	<NA>	gold	yellow
3	R2-D2	96	32	<NA>	white, blue	red

Special selection function

`dplyr` has special selection functions. See `?tidyselect::select_helpers`

- `contains()` Select columns that contain a character string
- `starts_with()` Select columns that start with a character string
- `ends_with()` Select columns that end with a string
- `matches()` Select columns that match a regular expression
- `everything()` Select all columns
- `num_range()` Select columns named something like x1, x2, x3, x4, x5
- `one_of(name_vector)` Select columns where the names are stored in a vector

Selection function examples

```
1 starwars %>%
2   select(name, ends_with("color")) %>% # selects name and cols endi
3   head(3)
```

A tibble: 3 × 4

	name	hair_color	skin_color	eye_color
	<chr>	<chr>	<chr>	<chr>
1	Luke Skywalker	blond	fair	blue
2	C-3PO	<NA>	gold	yellow
3	R2-D2	<NA>	white, blue	red

```
1 # selects name column and columns that match the regex, which says
2 starwars %>%
3   select(name, matches("s$")) %>%
4   head(3)
```

A tibble: 3 × 6

	name	mass	species	films	vehicles	starships
	<chr>	<dbl>	<chr>	<list>	<list>	<list>
1	Luke Skywalker	77	Human	<chr [5]>	<chr [2]>	<chr [2]>
2	C-3PO	75	Droid	<chr [6]>	<chr [0]>	<chr [0]>
3	R2-D2	32	Droid	<chr [7]>	<chr [0]>	<chr [0]>

Selecting with a variable

You can also select with a vector of names. To accomplish this, use the functions `all_of()` or `any_of()`

```
1 vars <- c("name", "mass", "height")
2 starwars %>% select(all_of(vars))
```

```
# A tibble: 87 × 3
```

	name	mass	height
	<chr>	<dbl>	<int>
1	Luke Skywalker	77	172
2	C-3PO	75	167
3	R2-D2	32	96
4	Darth Vader	136	202
5	Leia Organa	49	150
6	Owen Lars	120	178
7	Beru Whitesun Lars	75	165
8	R5-D4	32	97
9	Biggs Darklighter	84	183
10	Obi-Wan Kenobi	77	182

```
# i 77 more rows
```

Filter rows with `filter()`

With `filter()` you specify conditions to filter the rows in the data. Filter can use any condition that can be expressed as a logical vector with length equal to the number of rows.

```
1 starwars %>%
2   filter(name == "R2-D2")
```

```
# A tibble: 1 × 14
```

	name	height	mass	hair_color	skin_color	eye_color	birth_year	sex	gender
	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>	<chr>
1	R2-D2	96	32	<NA>	white	blue	red	33	none

```
masculine
```

```
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
```

```
#   vehicles <list>, starships <list>
```

filter() examples

Multiple conditions can be applied. Using the comma is equivalent to using &

```
1 starwars %>%  
2   filter(species %in% c("Human", "Droid"), height < 175)
```

```
# A tibble: 14 × 14
```

	name	height	mass	hair_color	skin_color	eye_color	birth_year	sex
gender	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>
	<chr>							
1	Luke Sk...	172	77	blond	fair	blue	19	male
mascu...								
2	C-3PO	167	75	<NA>	gold	yellow	112	none
mascu...								
3	R2-D2	96	32	<NA>	white, bl...	red	33	none
mascu...								
4	Leia Or...	150	49	brown	light	brown	19	fema...
femin...								
5	Beru Wh...	165	75	brown	light	blue	47	fema...
femin...								
6	D5-D4	97	32	<NA>	white	red	33	none

`filter()` is very powerful with regular expressions

We'll learn regular expressions in the next lecture.

`str_detect()` returns a logical vector.

```
1 starwars %>%
2   filter(str_detect(name, "^F")) # the name starts with F
```

A tibble: 2 × 14

	name	height	mass	hair_color	skin_color	eye_color	birth_year	sex
gender								
	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>
<chr>								
1	Finis Va...	170	NA	blond	fair	blue	91	male
mascu...								
2	Finn	NA	NA	black	dark	dark	NA	male
mascu...								

i 5 more variables: homeworld <chr>, species <chr>, films <list>,
vehicles <list>, starships <list>

The **dplyr** functions can be piped into each other

- use `|` for 'OR'

```
1 starwars %>%  
2   filter(hair_color == "none" | eye_color == "black") %>%  
3   select(name, species, homeworld, hair_color, eye_color)
```

A tibble: 39 × 5

	name	species	homeworld	hair_color	eye_color
	<chr>	<chr>	<chr>	<chr>	<chr>
1	Darth Vader	Human	Tatooine	none	yellow
2	Greedo	Rodian	Rodia	<NA>	black
3	IG-88	Droid	<NA>	none	red
4	Bossk	Trandoshan	Trandosha	none	red
5	Lobot	Human	Bespin	none	blue
6	Ackbar	Mon Calamari	Mon Cala	none	orange
7	Nien Nunb	Sullustan	Sullust	none	black
8	Nute Gunray	Neimodian	Cato Neimoidia	none	red
9	Jar Jar Binks	Gungan	Naboo	none	orange

Sort rows with `arrange()`

If you want to put things in descending order, wrap the variable name with `desc()`

```
1 starwars %>%  
2   select(name, birth_year, height, mass) %>%  
3   arrange(desc(birth_year), mass)
```

```
# A tibble: 87 × 4
```

	name	birth_year	height	mass
	<chr>	<dbl>	<int>	<dbl>
1	Yoda	896	66	17
2	Jabba Desilijic Tiure	600	175	1358
3	Chewbacca	200	228	112
4	C-3PO	112	167	75
5	Dooku	102	193	80
6	Ki-Adi-Mundi	92	198	82
7	Qui-Gon Jinn	92	193	89
8	Finis Valorum	91	170	NA
9	Palpatine	82	170	75
10	Cliegg Lars	82	183	NA

```
# i 77 more rows
```

Select rows based on their position with `slice()`

`slice()` lets you select rows based on their locations. The following selects rows 5 through 10

```
1 starwars %>% slice(5:10)
```

```
# A tibble: 6 × 14
```

	name	height	mass	hair_color	skin_color	eye_color	birth_year	sex
gender	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>
<chr>								
1	Leia Org...	150	49	brown	light	brown	19	fema...
	femin...							
2	Owen Lars	178	120	brown, gr...	light	blue	52	male
	mascu...							
3	Beru Whi...	165	75	brown	light	blue	47	fema...
	femin...							
4	R5-D4	97	32	<NA>	white, red	red	NA	none
	mascu...							
5	Biggs Da...	183	84	black	light	brown	24	male
	mascu...							

slice_sample()

`slice_sample()` lets you randomly select rows which can be useful to get a peek at portions of the entire tibble rather than just the head

```
1 starwars %>% slice_sample(n = 5)
```

```
# A tibble: 5 × 14
```

	name	height	mass	hair_color	skin_color	eye_color	birth_year	sex
gender	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>
<chr>								
1	Adi Gall...	184	50	none	dark	blue	NA	fema...
	femin...							
2	BB8	NA	NA	none	none	black	NA	none
	mascu...							
3	Mas Amed...	196	NA	none	blue	blue	NA	male
	mascu...							
4	Arvel Cr...	NA	NA	brown	fair	brown	NA	male
	mascu...							
5	Bib Fort...	180	NA	none	pale	pink	NA	male
	mascu...							

```
#> # A tibble: 5 × 14
```


slice_min() and slice_max()

`slice_min()` and `slice_max()` lets you select rows with the lowest or highest values in a variable. It is similar to using `arrange()` on a single variable and then `head()`.

```
1 starwars %>% slice_max(mass, n = 3)
```

```
# A tibble: 3 × 14
```

	name	height	mass	hair_color	skin_color	eye_color	birth_year	sex
gender	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>
	<chr>							
1	Jabba De...	175	1358	<NA>	green-tan...	orange	600	herm...
	mascu...							
2	Grievous	216	159	none	brown, wh...	green, y...	NA	male
	mascu...							
3	IG-88	200	140	none	metal	red	15	none
	mascu...							

```
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,  
#   vehicles <list>, starships <list>
```

Create new variables with `mutate()`

Use `mutate()` to create new variables based on existing variables. The new variable will be the last column, so we frequently use it with `select`.

```
1 starwars %>%
2   mutate(height_in = height / 2.54) %>% head(1)
```

```
# A tibble: 1 × 15
```

```
  name      height  mass hair_color skin_color eye_color birth_year sex
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr>
1 Luke Sky...  172    77 blond      fair      blue      19 male
mascu...
# i 6 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>, height_in <dbl>
```

```
1 starwars %>%
2   mutate(height_in = height / 2.54) %>%
3   select(name, height, height_in) %>% head(1)
```

```
# A tibble: 1 × 3
```

```
  name      height height_in
```

New variables must have the same number of rows

Important: Because `mutate()` adds a new column to the data set, the variable you are creating must have the same number of values as rows in the data set.

```
1 starwars %>%  
2   select(name, mass) %>%  
3   mutate(cumulative_mean = cummean(mass))
```

A tibble: 87 × 3

	name <chr>	mass <dbl>	cumulative_mean <dbl>
1	Luke Skywalker	77	77
2	C-3PO	75	76
3	R2-D2	32	61.3
4	Darth Vader	136	80
5	Leia Organa	49	73.8
6	Owen Lars	120	81.5
7	Beru Whitesun Lars	75	80.6
8	R5-D4	32	74.5

9	Biggs Darklighter	84	75.6
10	Obi-Wan Kenobi	77	75.7

77 more rows

Some useful functions for `mutate()`

- `pmin()`, `pmax()` Element-wise min and max
- `cummin()`, `cummax()` Cumulative min and max
- `cumsum()`, `cumprod()` Cumulative sum and product
- `between()` Are values between a and b?
- `cummean()` Cumulative mean
- `lead()`, `lag()` Copy values with offset
- `ntile()` Bin vector into n buckets

mutate() examples

```
1 starwars %>%
2   select(name, mass, birth_year) %>%
3   mutate(
4     cummin_mass = cummin(mass), # cummin gives the min value seen s
5     ratio = mass / mean(mass, na.rm = TRUE), # we divide mass/by th
6     massyear_pmin = pmin(mass, birth_year), # pmin gives the elemer
7     lag2 = lag(massyear_pmin, 2)) # lag offsets the column values
```

A tibble: 87 × 7

	name <chr>	mass <dbl>	birth_year <dbl>	cummin_mass <dbl>	ratio <dbl>	massyear_pmin <dbl>	lag2 <dbl>
1	Luke Skywalker	77	19	77	0.791	19	NA
2	C-3PO	75	112	75	0.771	75	NA
3	R2-D2	32	33	32	0.329	32	19
4	Darth Vader	136	41.9	32	1.40	41.9	75
5	Leia Organa	49	19	32	0.504	19	32
6	Owen Lars	120	52	32	1.23	52	41.9
7	Beru Whitesun Lars	75	47	32	0.771	47	19
8	R5-D4	32	NA	32	0.329	NA	52
9	Biggs Darklighter	84	24	32	0.863	24	47
10	Obi-Wan Kenobi	77	57	32	0.791	57	NA

i 77 more rows