

# Video 24: regex character sets

Stats 102A

Miles Chen, PhD

# Regular Expressions

# Resources

## Cheat Sheets for `stringr` and Regular Expressions

- Strings Cheat Sheet at: <https://posit.cloud/learn/cheat-sheets>
- <https://www.cheatography.com//davechild/cheat-sheets/regular-expressions/pdf/>

## Sites for Testing Regular Expressions

- <https://regex101.com/>
- <https://regexr.com/>

# Regular Expressions

One main application of string manipulation is pattern matching. Finding patterns in text are useful for data validation, data scraping, text parsing, filtering search results, etc.

A **regular expression** (or **regex**) is a set of symbols that describes a text pattern. More formally, a regular expression is a pattern that describes a set of strings.

Regular expressions are a formal language in the sense that the symbols have a defined set of rules to specify the desired patterns. The best way to learn the syntax and become fluent with regular expressions is to practice.

# Applications of Regular Expressions

Some common applications of regular expressions:

- Test if a phone number has the correct number of digits
- Test if a date follows a specific format (e.g. mm/dd/yy)
- Test if an email address is in a valid format
- Test if a password has numbers and special characters
- Search a document for gray spelled either as “gray” or “grey”
- Search a document and replace all occurrences of “Will”, “Bill”, or “W.” with “William”
- Count the number of times in a document that the word “analysis” is immediately preceded by the words “data”, “computer”, or “statistical”
- Convert a comma-delimited file into a tab-delimited file
- Find duplicate words in a text

# stringr for Regular Expressions

R has native regex handling capabilities (e.g. `grep()`), but `stringr` has made their usage easier and more consistent.

Function	Description
<code>str_detect(str, pattern)</code>	Detect the presence of a pattern and returns <code>TRUE</code> if it is found
<code>str_locate(str, pattern)</code>	Locate the 1st position of a pattern and return a matrix with start & end.
<code>str_extract(str, pattern)</code>	Extracts text corresponding to the first match.
<code>str_match(str, pattern)</code>	Extracts capture groups formed by <code>()</code> from the first match.
<code>str_split(str, pattern)</code>	Splits string into pieces and returns a list of character vectors.

# Literal Characters

The most basic type of regular expressions are **literal characters**, which are characters that match themselves.

A literal character match is one in which a given character such as the letter "R" matches the letter R. This type of match is the most basic type of regular expression operation: just matching plain text.

All the letters and digits in the English alphabet (i.e., alphanumeric characters) are considered literal characters because, as regular expressions, they match themselves.

# Matching Literal Characters

Literal character matching is case sensitive.

```
1 str_locate("I love stats", "stat")
```

```
      start end  
[1,]      8  11
```

```
1 str_locate("I love Stats", "stat")
```

```
      start end  
[1,]     NA  NA
```



# Matching Literal Characters

The `str_locate()` function only returns the first occurrence of a match. To find all matches, use `str_locate_all()`.

```
1 love_stats <- "I love statistics, so I am a stats major."
```

```
1 str_locate(love_stats, "stat")
```

```
      start end  
[1,]      8  11
```

```
1 str_locate_all(love_stats, "stat")
```

```
[[1]]  
      start end  
[1,]      8  11  
[2,]     30  33
```

# Metacharacters

Not all characters match themselves. Any character that is not a literal character is a **metacharacter**.

The power of regular expressions comes from the ability to use a number of special metacharacters that modify how the pattern matching is performed.

The list of metacharacters used in regular expressions is given below:

```
1 . ^ $ * + ? { } [ ] \ | ( )
```

# The Wild Metacharacter

The **dot** (or **period**) `.` is called the **wild** metacharacter (sometimes the **wildcard**). This metacharacter is used to match ANY (single) character except a new line.

```
1 not <- c("not", "note", "knot", "nut")
```

```
1 str_detect(not, "n.t")
```

```
[1] TRUE TRUE TRUE TRUE
```

```
1 str_detect(not, "no.")
```

```
[1] TRUE TRUE TRUE FALSE
```

**Question:** Consider the following vector.

```
1 fives <- c("5.00", "5100", "5-00", "5 00")
```

What will `str_detect(fives, "5.00")` return?

# Escaping Metacharacters

Because of their special properties, metacharacters cannot be matched directly as literal characters.

To do a literal match, we need to **escape** the metacharacter by adding a backslash `\` in front of the metacharacter.

In R, however, since the backslash `\` itself is a metacharacter for normal strings, we need a **double backslash** `\\` to escape metacharacters in regular expressions.

```
1 str_detect("abc[def", "[")
```

```
## Error in stri_detect_regex(string, pattern, negate = negate,  
opts_regex = opts(pattern)) : Missing closing bracket on a bracket  
expression. (U_REGEX_MISSING_CLOSE_BRACKET)
```

```
1 str_detect("abc[def", "\\[")
```

```
[1] TRUE
```

# Escaping The Wild Metacharacter

Consider the following vector.

```
1 fives <- c("5.00", "5100", "5-00", "5 00")  
2 str_detect(fives, "5.00")
```

```
[1] TRUE TRUE TRUE TRUE
```

To match the literal dot `.` and match only `5.00`, we need to escape the wild metacharacter:

```
1 str_detect(fives, "5\\.00")
```

```
[1] TRUE FALSE FALSE FALSE
```

# Character Sets

Square brackets `[ ]` indicate a **character set**, which will match any one of the characters that are inside the set.

A character set will match only one character. The order of the characters inside the set does not matter.

For example, the character set defined by `[aeiou]` will match any one lower case vowel.

```
1 pnx <-  
2   c("pan", "pen", "pin", "p0n", "p.n", "paun", "pwn3d", "happiness")
```

```
1 str_detect(pnx, "p[aeiou]n")
```

```
[1]  TRUE  TRUE  TRUE FALSE FALSE FALSE FALSE  TRUE
```

# Character Ranges

To match all capital letters (in English), we can define the character set `[ABCDEFGHIJKLMNOPQRSTUVWXYZ]`.

However, this expression is long and inconvenient. Fortunately, the **dash** `-` metacharacter is a shortcut to indicate a **range** of characters.

Some examples:

Pattern	Character Range
<code>[a-q]</code>	All lower case letters from a to q
<code>[A-Q]</code>	All upper case letters from A to Q
<code>[a-zA-Z]</code>	All 52 ASCII letters
<code>[0-7]</code>	All digits from 0 to 7

# Character Ranges

Character ranges are useful to match various occurrences of a certain type of character.

**Question:** Consider the following vector.

```
1 triplets <- c("123", "abc", "ABC", ":-)", "ab12a", "a8908ab")
```

What pattern can be defined to match 3 adjacent digits?

```
1 str_detect(triplets, "[0-9][0-9][0-9]")
```

```
[1] TRUE FALSE FALSE FALSE FALSE TRUE
```

Similarly, the pattern `[a-z][a-z][a-z]` matches three adjacent lower case letters.



# Negative Character Sets

A common situation when working with regular expressions consists of matching characters that are NOT part of a certain set.

This type of matching can be done using a **negative character set**: by matching any one character that is not in the set.

The **caret** `^` metacharacter is used to create negative character sets.

If a caret `^` is placed in the first position inside a character set, it means **negation** (similar to the negative sign in numeric indices or the exclamation point in logical expressions).

For example, the pattern `[^aeiou]` means “not any one of lower case vowels.”

**Note:** The caret `^` is a metacharacter that has more than one meaning depending on where it appears in a pattern.

# Negative Character Sets

For example, the pattern `[^A-Z]` will match any character that is NOT an upper case letter.

```
1 basic <- c("1", "a", "A", "&", "-", "^")
```

```
1 str_detect(basic, "[^A-Z]")
```

```
[1] TRUE TRUE FALSE TRUE TRUE TRUE
```

# Negative Character Sets

**Caution:** It is important that the caret `^` is the first character inside the character set, otherwise the set is not a negative one.

For example, the pattern `[A-Z^]` mean any one upper case letter or the caret character.

```
1 basic
```

```
[1] "1" "a" "A" "&" "-" "^"
```

```
1 str_detect(basic, "[A-Z^]")
```

```
[1] FALSE FALSE TRUE FALSE FALSE TRUE
```

**Question:** What pattern means “anything except the caret?”

# Negative Character Sets

The pattern `[^\^]` can be used to mean anything except the caret.

```
1 basic
```

```
[1] "1" "a" "A" "&" "-" "^"
```

```
1 str_detect(basic, "[^\^]")
```

```
[1] TRUE TRUE TRUE TRUE TRUE FALSE
```

# Metacharacters Inside Character Sets

Most metacharacters inside a character set are already escaped. This implies that you do not need to escape them using double backslashes.

```
1 pnx <-  
2   c("pan", "pen", "pin", "p0n", "p.n", "paun", "pwn3d")
```

```
1 str_detect(pnx, "p[ae.iou]n")
```

```
[1]  TRUE  TRUE  TRUE FALSE  TRUE FALSE FALSE
```

The dot `.` inside the character set now represents the literal dot character rather than the wildcard character.

**Note:** Not all metacharacters become literal characters when they appear inside a character set. The exceptions are the opening bracket `[`, the closing bracket `]`, the dash `-`, the caret `^` (if at the front or by itself), and the backslash `\`.

# Character Classes

Closely related to character sets and character ranges are **character classes**, which are used to match a certain class of characters.

The most common character classes in most regex engines are:

Pattern	Matches	Same as
<code>\\d</code>	Any digit	<code>[0-9]</code>
<code>\\D</code>	Any non-digit	<code>[^0-9]</code>
<code>\\w</code>	Any word character	<code>[a-zA-Z0-9_]</code>
<code>\\W</code>	Any non-word character	<code>[^a-zA-Z0-9_]</code>
<code>\\s</code>	Any whitespace character	<code>[\f\n\r\t\v]</code>
<code>\\S</code>	Any non-whitespace character	<code>[^\f\n\r\t\v]</code>

Character classes can be thought of as another type of metacharacter or as shortcuts for special character sets.

# Whitespace

There are several types of whitespace characters, shown in the following table:

Character	Description
<code>\f</code>	Form feed (page break)
<code>\n</code>	Line feed (new line)
<code>\r</code>	Carriage return
<code>\t</code>	Tab
<code>\v</code>	Vertical tab

For situations with non-printing whitespace characters, it can be difficult to determine which exact character it is, so the whitespace class `\s` is a useful way to match with all of them.

# Character Classes

For example:

```
1 pnx <-  
2   c("pan", "pen", "pin", "p0n", "p.n", "paun", "pwn3d")
```

```
1 str_detect(pnx, "p\\d") # p followed by digit
```

```
[1] FALSE FALSE FALSE  TRUE FALSE FALSE FALSE
```

```
1 str_detect(pnx, "p\\D") # p followed by non-digit
```

```
[1]  TRUE  TRUE  TRUE FALSE  TRUE  TRUE  TRUE
```

```
1 str_detect(pnx, "p\\W") # p followed by non-word character
```

```
[1] FALSE FALSE FALSE FALSE  TRUE FALSE FALSE
```



# POSIX Character Classes

There is another type of character classes known as **POSIX character classes** that is supported by the regex engine in R.

Class	Description	Same as
<code>[ :alnum: ]</code>	Any letter or digit	<code>[ a-zA-Z0-9 ]</code>
<code>[ :alpha: ]</code>	Any letter	<code>[ a-zA-Z ]</code>
<code>[ :digit: ]</code>	Any digit	<code>[ 0-9 ]</code>
<code>[ :lower: ]</code>	Any lower case letter	<code>[ a-z ]</code>
<code>[ :upper: ]</code>	Any upper case letter	<code>[ A-Z ]</code>
<code>[ :space: ]</code>	Any whitespace, including space	<code>[ \f\n\r\t\v ]</code>
<code>[ :punct: ]</code>	Any punctuation symbol	
<code>[ :print: ]</code>	Any printable character	
<code>[ :graph: ]</code>	Any printable character excluding space	
<code>[ :xdigit: ]</code>	Any hexadecimal digit	<code>[ a-fA-F0-9 ]</code>

# POSIX Character Classes

To use POSIX classes in R, the class needs to be wrapped inside a regex character class, i.e., the class needs to be inside a second set of square brackets.

For example:

```
1 pnx <-  
2   c("pan", "pen", "pin", "p0n", "p.n", "paun", "pwn3d")  
3 str_detect(pnx, "[[:alpha:]]") # has any letter
```

```
[1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

```
1 str_detect(pnx, "[[:digit:]]") # has any digit
```

```
[1] FALSE FALSE FALSE  TRUE FALSE FALSE  TRUE
```

